

DP Syntropy → Marketing Mix Modeling: Status & What We Need to Go Live

In one line: The pipeline is now instrumented so that Marketing Mix Modeling becomes possible the moment real campaigns run — the measurement machine builds its own training data automatically. The only thing left to build is the model itself, and that is waiting on real traffic, not on engineering.

The idea (why MMM)

Marketing exists to capture revenue. The way to know whether our outreach actually captures it — rather than just staying busy — is Marketing Mix Modeling: use regression to measure **which channels actually move a specific objective**, so we optimize the mix instead of "doing more." For that to work, three things must be true, and we have now built all three into the pipeline:

1. **Every campaign pursues one limited, measurable objective** — not "engagement," but a concrete target (e.g. *one reply-or-meeting within 21 days*).
2. **The objective is tied to a funnel stage** — cold, aware, engaged, qualified, or customer — so we can learn what works *per stage*.
3. **Every outreach touch and its outcome are recorded** in a form a regression can read.

What is now automated

- All 18 campaign playbooks now carry a **limited objective** (a success metric, a numeric target, and a time window) and a **funnel stage**.
- When a campaign is created, the system **automatically creates that campaign's objective** — no manual step.
- A daily job **automatically marks each objective met or missed** within its window, by checking whether the right success signal actually happened.
- Each account's campaign produces one row in a "**marketing-mix matrix**": the channels used (email, LinkedIn, phone, ads, webinar, etc.) alongside whether the objective was hit.

That matrix is the training data for MMM, and **it fills itself** as campaigns run. No further wiring is required.

The one-account picture

An account enters, lands at a funnel stage, and is given that stage's limited objective. Reps work the touches; the system records the mix. Within the window, the objective auto-resolves to met or missed. Hit it, and the account graduates to the next stage with a new objective. Every such cycle adds one clean, labeled row to the matrix.

What is left

Only the model. Once we have enough real campaigns, we pull the matrix into a regression and get, per funnel stage, the evidence of which channels drive the objective — which then feeds our channel recommendation engine so it optimizes on evidence rather than heuristics. This is the "nirvana," and it is

deliberately the last step because a regression needs data volume and variety that a prototype does not yet have.

What we need from the company to light it up

Blockers — to prove the loop on one real account:

- A **live marketing website + landing pages** (real destinations for tracked links; today they are placeholders, which is why no results are being measured yet).
- The **live WHEN-engine endpoint** and a real feed of account briefs.
- **Real target accounts** with contact emails.
- A **designated rep** to execute the outreach and log outcomes.

Scale — to reach MMM:

- **Volume**: a sustained flow of real accounts (regression needs many observations).
- **Sign-off on the real objectives** per funnel stage (what success actually is for us).
- **Company-owned accounts** (the platform currently runs on personal accounts) and a decision on the CRM environment.

Bottom line

The measurement-and-learning half of the pipeline — the part that turns marketing into a self-improving system — is **built, tested, and dormant, waiting only for real traffic**. The next move is not more engineering; it is securing the inputs above so the loop can start turning on a first real account.

Prepared 2026-07-07. Prototype/demo status — not yet production.